Lecture 06: Data Manipulation

Taylor Arnold

Data manipulation

- ▶ so far we have primarily been working with a single dataset as it was directly given to us. In many cases it is advantageous to produce new datasets from our original one.

- ▶ can be as simple as selecting a subset of the original columns or rows, or as involved as changing the *level of analysis* of the entire dataset.

- ▶ we will introduce several functions called *verbs* that assist in manipulating datasets. These all come from the package **dplyr**, supplemented with functions in **smodels**.

select

In this handout, as with the others so far, I will use the `msleep` dataset in order to show various numerical summaries. Rather than the entire dataset, it will be helpful to take just a subset of the columns. This gives us the first verb we'll see: `select`:

```
msleep <- select(msleep, awake, vore, genus)
```

A grammar for data manipulation

The filter function takes a dataset followed by a logical statements. It returns a dataset that has any rows in the input data that match the filtering statements. For example, the following returns a dataset for all mammals that are awake more than 20 hours per day:

```
filter(msleep, awake > 20)
```

Filtering rows

```
filter(msleep, awake > 20)
```

```
## # A tibble: 9 x 3
##   awake vore      genus
##   <dbl> <chr>     <chr>
## 1 21.00 herbi     Capreolus
## 2 20.10 herbi     Elephas
## 3 21.10 herbi     Equus
## 4 20.90 herbi     Equus
## 5 22.10 herbi     Giraffa
## 6 21.35 carni Globicephalus
## # ... with 3 more rows
```

Filtering rows

It is very important to notice that the original `msleep` dataset has
not been altered here. If we want to actually work with the filtered
data, we need to save it using the assignment operator `<-` and
given it a name. To pull out just the *sleepy* mammals, we could do
the following:

```
sleepy <- filter(msleep, awake < 6)
```

You should see now that a new dataset appears in your workspace named sleepy. It is possible to work with this new dataset in all of the ways we have plotted and (now) filtered the original data.

Constructing new variables

The mutate function preserves all rows of the original dataset, unlike filter, but adds a new variables. For example, to add hours asleep into the dataset we can do this:

```
msleep <- mutate(msleep, asleep = 24 - awake)
```

Here, I used the assignment command to save the result back into the msleep dataset.

If you look in your environment window, you'll still see just a single version of `msleep` but this one will have one extra variables. This is a relatively safe pratice with the mutate function, as we are simply adding information, but should be generally avoided when using filter.

It is possible to use mutate to redefine an existing variable by giving mutate a variable name that already exists. Be careful of this, particularly if you are overwritting the original dataset.

Constructing new variables

Of course, we could also do the same thing with the following:

```
msleep$asleep <- 24 - msleep$awake
```

It is just nice to use the `mutate` function as it is consistent with all of the other **dplyr** verbs.

Summarizing data

The group_summarize function is, in my opinion, the most complex verb that we will use this semester. If we use the function on a dataset without any other options it gives the mean, median, standard deviation, and sum for every numeric variable in the dataset. An overall count is also included. Let's apply it to msleep:

```
group_summarize(msleep)
```

```
group_summarize(msleep)
```

```
## # A tibble: 1 x 9
##   awake_mean asleep_mean awake_median asleep_median
##        <dbl>       <dbl>        <dbl>         <dbl>
## 1   13.56747    10.43253         13.9          10.1
## # ... with 5 more variables: awake_sd <dbl>,
## #  asleep_sd <dbl>, awake_sum <dbl>, asleep_sum <dbl>,
## #   n <int>
```

Summarizing data

These variables could have easily been computed by calling the respective functions individually in R. The group summarize function becomes more interesting when we pass it a second input giving a variable to group by. For example, here is the summary *grouped by vore*

```
group_summarize(msleep, vore)
```

Summarizing data

```
group_summarize(msleep, vore)

## # A tibble: 5 x 10
##     vore awake_mean asleep_mean awake_median asleep_median
##     <chr>     <dbl>      <dbl>        <dbl>         <dbl>
## 1  carni  13.62632   10.373684        13.6          10.4
## 2  herbi  14.49062    9.509375        13.7          10.3
## 3 insecti  9.06000   14.940000         5.9          18.1
## 4   omni  13.07500   10.925000        14.1           9.9
## 5   <NA>  13.81429   10.185714        13.4          10.6
## # ... with 5 more variables: awake_sd <dbl>,
## #  asleep_sd <dbl>, awake_sum <dbl>, asleep_sum <dbl>,
## #   n <int>
```

Summarizing data

Notice that the result now provides these summaries for each group. It is possible to summarize by multiple groups at once, which produces summaries for each unique combination of the those variables. For instance, we could summarize by both genus and vore:

```
group_summarize(msleep, vore, genus)
```

Which returns a row for each unique combination of vore and genus.

Summarizing data

```
group_summarize(msleep, vore, genus)

## # A tibble: 77 x 11
##   vore      genus awake_mean asleep_mean awake_median
##   <chr>     <chr>      <dbl>       <dbl>        <dbl>
## 1 carni   Acinonyx      11.9        12.1         11.9
## 2 carni Callorhinus     15.3         8.7         15.3
## 3 carni      Canis      13.9        10.1         13.9
## 4 carni    Dasypus       6.6        17.4          6.6
## 5 carni      Felis      11.5        12.5         11.5
## 6 carni    Genetta      17.7         6.3         17.7
## # ... with 71 more rows, and 6 more variables:
## #  asleep_median <dbl>, awake_sd <dbl>, asleep_sd <dbl>,
## #   awake_sum <dbl>, asleep_sum <dbl>, n <int>
```

The summarized data also includes a column called n giving the total number of observations within a group. The count function functions similarly, but does not returns any of the other summary functions. This is convenient if all you need are counts:

```
count(msleep, vore)
```

Counting data

```
count(msleep, vore)

## # A tibble: 5 x 2
##      vore     n
##     <chr> <int>
## 1   carni    19
## 2   herbi    32
## 3 insecti     5
## 4    omni    20
## 5    <NA>     7
```

The sort option gives us the ability to order to data at the same time as it is counted:

```
count(msleep, vore, sort = TRUE)
```

Counting data

```
count(msleep, vore, sort = TRUE)

## # A tibble: 5 x 2
##      vore      n
##     <chr> <int>
## 1   herbi    32
## 2    omni    20
## 3   carni    19
## 4    <NA>     7
## 5 insecti     5
```

Combining datasets

The final verb that we will use this semester is also the only two-table verb that we will need. It will be used to combine a dataset with metadata about one or more of its variables. To illustrate, let's make a small data frame that contains the full name for the short-hand abbreviations given in the variable `vore`.

```
meta <- data_frame(
          vore = c("carni", "omni", "herbi", "insecti"),
       full_name = c("carnivore", "omnivore",
                       "herbivore", "insectivore"))
```

left_join

```
meta

## # A tibble: 4 x 2
##      vore  full_name
##     <chr>      <chr>
## 1   carni  carnivore
## 2    omni   omnivore
## 3   herbi  herbivore
## 4 insecti insectivore
```

left_join

To combine these with the original dataset, we use the left_join function, giving the larger dataset first:

```
msleep <- left_join(msleep, meta)
```

There is a new variable full_name that now displays the full name for each vore type.

```
msleep
```

```
## # A tibble: 83 x 5
##    awake vore       genus asleep full_name
##    <dbl> <chr>      <chr>  <dbl>     <chr>
## 1  11.9 carni    Acinonyx   12.1 carnivore
## 2   7.0 omni        Aotus   17.0 omnivore
## 3   9.6 herbi   Aplodontia   14.4 herbivore
## 4   9.1 omni       Blarina   14.9 omnivore
## 5  20.0 herbi         Bos    4.0 herbivore
## 6   9.6 herbi     Bradypus   14.4 herbivore
## # ... with 77 more rows
```