


## Lecture 19: Scraping Wikipedia Lists

---

Taylor Arnold



Today, we are going to use R to systemically build a dataset from information on Wikipedia.

Words of caution about scraping data (note: this is not legal advice):

- ▶ many websites prohibit the automatic downloading of material in their terms and conditions
- ▶ a small set of people have gotten in serious trouble for scraping websites against their terms of use
- ▶ we are only going to apply this to Wikipedia, as they have generously permissive rules about this sort of thing
- ▶ if you ever want to replicate this on a large scale, please use the Wikipedia API rather than the tools here

Wikipedia has curated a table of the most populous cities in the world. We can download this list, as in the last class, using the following code:

```
url <- "https://en.wikipedia.org/wiki/List_of_largest_cities"  
wpage <- data_frame(line = readLines(url))
```

Make sure to view the site itself as well in a browser.

Searching through the source HTML, we see that rows in the city table always start with the string `<th scope="row">`. Let's filter our dataset to only include these rows:

```
wpage <- filter(wpage, stri_detect(line, fixed = '<th scope="row">'))
```

## extracting city names

We will also remove all HTML tags on these rows, and store the results as the variable `city`.

```
wpage <- mutate(wpage, city = stri_replace_all(line, "", regex = "<[>]"))
wpage$city
```

```
##      [1] "Shanghai"      "Karachi"
##      [3] "Beijing"        "Delhi"
##      [5] "Lagos"          "Dhaka"
##      [7] "Guangzhou"      "Istanbul"
##      [9] "Tokyo"          "Mumbai"
##     [11] "Moscow"         "São Paulo"
##     [13] "Shenzhen"       "Suzhou"
##     [15] "Lahore"         "Cairo"
##     [17] "Kinshasa"       "Jakarta"
##     [19] "Seoul"          "Mexico City"
##     [21] "Lima"           "London"
##     [23] "Xi'an"          "New York City"
##     [25] "Bengaluru"      "Bangkok"
##     [27] "Nanjing"        "Dongguan"
```

Finally, looking at row of data, notice that there is a link to the Wikipedia page about each city:

```
wpage$line[1]
```

```
## [1] "<th scope=\"row\"><a href=\"/wiki/Shanghai\" title=\"Shanghai\">S
```

## extracting links

We can extract these using the `stri_extract` function:

```
wpage <- mutate(wpage, link = stri_extract(line, regex = "/wiki/[^\"]"])  
wpage$link
```

```
##      [1] "/wiki/Shanghai"  
##      [2] "/wiki/Karachi"  
##      [3] "/wiki/Beijing"  
##      [4] "/wiki/Delhi"  
##      [5] "/wiki/Lagos"  
##      [6] "/wiki/Dhaka"  
##      [7] "/wiki/Guangzhou"  
##      [8] "/wiki/Istanbul"  
##      [9] "/wiki/Tokyo"  
##     [10] "/wiki/Mumbai"  
##     [11] "/wiki/Moscow"  
##     [12] "/wiki/S%C3%A3o_Paulo"  
##     [13] "/wiki/Shenzhen"  
##     [14] "/wiki/Suzhou"  
##     [15] "/wiki/Lahore"
```



With this link, we can now download the specific information from a given city. For instance, let's take the 5'th city (Lagos) and download the website from Wikipedia:

```
i <- 5
url <- stri_c("https://en.wikipedia.org", wpage$link[i])
cpage <- data_frame(line = readLines(url))
```

Note the use of the `stri_c` function to combine the base URL with the link.

One piece of information on most city pages is the latitude and longitude of the city. We can find this by detecting the string "Coordinates:" in the webpage. Here, we save only the first occurrence of the string (here there is only one, but we want to safely abstract this to other pages):

```
coord <- filter(cpage, stri_detect(line,  
                                fixed = "Coordinates:"))$line[1]
```

```
coord
```

```
## [1] "<td colspan=\"2\" style=\"text-align:center\">Coordinates: <span
```

Looking at the output, we want to not just remove the HTML tags but to capture a very specific part of the tags. Here we get the easiest form of the coordinates from the “geo” tag:

```
coord <- stri_extract(coord,  
                      regex = "<span class=\"geo\">[~<]+</span>")  
coord
```

```
## [1] "<span class=\"geo\">6.455027; 3.384082</span>"
```

We want to now remove the html tags. We could do this by a regular expression, but let's instead do it by taking a substring:

```
coord <- stri_sub(coord, 19, -8)  
coord
```

```
## [1] "6.455027; 3.384082"
```

Next, we want to split the string into two parts based on the “;” symbol. We do this with the `stri_split` function.

```
coord <- stri_split(coord, fixed = ";")[[1]]  
coord
```

```
## [1] "6.455027" " 3.384082"
```

Finally, we want to convert these strings into numbers. Recall that `factor` turns numeric data into categorical ones; at the time, I mentioned that rarely can we go in the reverse order. Here we use the `as.numeric` function to do just that:

```
coord <- as.numeric(coord)
lat <- coord[1]
lon <- coord[2]
lat
```

```
## [1] 6.455027
```

```
lon
```

```
## [1] 3.384082
```

Let's also find the name of the country that each city is in. this is a bit complicated because the row that shows the tag "Country" is actually one away from the row that gives the data:

```
filter(cpage, stri_detect(line,  
  fixed = "<th scope=\"row\">Country</th>"))
```

```
## # A tibble: 1 x 1  
##               line  
##             <chr>  
## 1 "<th scope=\"row\">Country</th>"
```

To fix this, wrap the `stri_detect` function in the function `which`. This gives the indices where something is `True`:

```
id <- which(stri_detect(cpage$line,  
                      fixed = "<th scope=\"row\">Country</th>"))  
id
```

```
## [1] 104
```



Now, we just add 1 to this index (extracting just the first one, in case there are multiple) and grab those lines:

```
country <- cpage$line[id[1] + 1]  
country
```

```
## [1] "<td><span class=\"flagicon\"><img alt=\"\" src=\"//upload.wikimed
```

Cleaning up the results yields the desired information:

```
country <- cpage$line[which(stri_detect(cpage$line,  
    fixed = "<th scope=\"row\">Country</th>"))[1] + 1]  
country <- stri_replace_all(country, "", regex = "<[^>]+>")  
country <- stri_replace_all(country, "", fixed = "&#160;")  
country
```

```
## [1] "Nigeria"
```

We now have code that works on a single city page. Let's add empty attributes to the `wpage` dataset. We'll fill these in in a moment.

```
wpage$lat <- NA  
wpage$lon <- NA  
wpage$country <- NA
```


We now use a *for loop* to cycle over all rows of the `wpage` dataset as follows:


```
for (i in 1:nrow(wpage)) {  
  url <- stri_c("https://en.wikipedia.org", wpage$link[i])  
  cpage <- data_frame(line = readLines(url))  
  
  # extract stuff  
  
  wpage$lat[i] <- lat  
  wpage$lon[i] <- lon  
  wpage$country[i] <- country  
}
```

```
for (i in 1:nrow(wpage)) {  
  url <- stri_c("https://en.wikipedia.org", wpage$link[i])  
  cpage <- data_frame(line = readLines(url))  
  
  coord <- filter(cpage,  
    stri_detect(line, fixed = "Coordinates:"))$line[1]  
  coord <- stri_extract(coord,  
    regex = "<span class=\"geo\">[^<]+</span>")  
  coord <- stri_sub(coord, 19, -8)  
  coord <- stri_split(coord, fixed = ";")[[1]]  
  coord <- as.numeric(coord)  
  wpage$lat[i] <- coord[1]  
  wpage$lon[i] <- coord[2]  
  
  country <- cpage$line[which(stri_detect(cpage$line,  
    fixed = "<th scope=\"row\">Country</th>"))[1] + 1]  
  country <- stri_replace_all(country, "", regex = "<[^>]+>")  
  country <- stri_replace_all(country, "", fixed = "&#160;")
```

We can now plot the data as follows:

```
qplot(lon, lat, data = wpage, label = city, geom = "text") +  
  theme_minimal()
```







Or tabulate by country:

```
count(wpage, country, sort = TRUE)
```

```
## # A tibble: 69 x 2
```

```
##           country      n
##           <chr> <int>
## 1           <NA>    69
## 2          China    21
## 3          India    16
## 4 People's Republic of China 10
## 5          Pakistan    9
## 6          Russia     7
## # ... with 63 more rows
```